

Technical Companion and Experiment Record

Post-Training Qwen2.5-0.5B with SFT, Reward Modeling, PPO, and DPO

Dheeraj Dhillon

July 2026

Abstract

This companion preserves the engineering, experimental, and artifact-level detail behind the shorter two-column project report. It is intentionally more operational: it records the custom-to-TRL migration, implementation checks, historical runs, Colab persistence, checkpoint semantics, metric definitions, artifact paths, and the subsequent DPO experiment design. The concise project report remains the primary academic account; this document is the durable technical record.

This report presents an end-to-end reinforcement learning from human feedback (RLHF) study on QWEN2.5-0.5B-INSTRUCT using NVIDIA’s HELPSTEER3-PREFERENCE. The project began with a custom implementation of response-only supervised fine-tuning (SFT), pairwise reward modeling, token-level Proximal Policy Optimization (PPO), KL control, generalized advantage estimation, checkpointing, and policy evaluation. After the custom system exposed practical failure modes, the training loops were migrated to Hugging Face TRL while repository-owned preprocessing, experiment tracking, evaluation, repetition diagnostics, and qualitative curation were retained.

The final pipeline trains on 36,264 filtered preference records with a 4,096-token sequence budget. SFT obtains a validation mean token accuracy of 72.02%. A two-epoch reward model obtains 65.62% pairwise validation accuracy on 1,917 preference pairs. The selected PPO policy is evaluated after 6,400 rollout episodes and 100 optimizer updates. On 2,017 validation prompts, the learned reward model prefers PPO over the original instruction model in 1,027 cases (50.92%), prefers the base model in 981 cases, and assigns 9 ties. The mean PPO-minus-Base reward delta is +0.6497.

That headline is not treated as proof of general model improvement. PPO responses are longer, reach the 1,024-token evaluation cap in 27.42% of cases versus 8.82% for Base, and exceed a 25% repeated 4-gram threshold in 31.88% of cases versus 10.11% for Base. The project therefore contributes a reproducible alignment system and a transparent evaluation case study rather than a claim that PPO universally improves a 0.5B model. All 2,017 prompt and response comparisons are exposed through a public response explorer, with a balanced 100-example subset for faster review.

Contents

1	Introduction	4
1.1	Contributions	4
2	RLHF Formulation	4
2.1	Notation	4
2.2	Supervised fine-tuning	5
2.3	Pairwise reward modeling	5
2.4	KL-regularized PPO	5
3	Model, Data, and Preprocessing	6
3.1	Base model and preference data	6
3.2	Chat formatting and truncation	7

4	Implementation and Experimental Infrastructure	7
4.1	From a custom trainer to TRL	7
4.2	Direct preference optimization extension	8
4.3	Implementation details adopted from N+	8
4.4	Reproducibility and Colab execution	9
5	Final Training Configuration and Results	9
5.1	Supervised fine-tuning	9
5.2	Reward-model training	10
5.3	PPO alignment	11
6	Experimental Progression	11
7	Final Policy-Suite Evaluation	12
7.1	Protocol	12
7.2	Overall results	13
7.3	Domain behavior	14
8	Qualitative Audit and Response Explorer	14
8.1	Stopping and repetition	14
8.2	Deterministic triage	15
8.3	Curated 100 and complete public artifact	16
9	Discussion	16
9.1	What the final run demonstrates	16
9.2	Limitations	16
10	Future Work	17
10.1	Establish a controlled evaluation protocol	17
10.2	Improve reward-model reliability	17
10.3	Strengthen supervised fine-tuning	18
10.4	Redesign PPO around observed failures	18
10.5	Compare alternative preference objectives	19
10.6	Improve decoding and stopping	19
10.7	Add factuality, execution, and safety evaluators	19
10.8	Scale model and systems capacity deliberately	19
10.9	Reproducibility and artifact governance	20
10.10	Practical next sequence	20
11	Conclusion	20

A Detailed Historical Experiment Ledger	20
A.1 Initial validation and sequence-length diagnosis	20
A.2 Frozen custom SFT, reward, and PPO baseline	21
A.3 Historical policy-suite results	21
A.4 Historical A100 40GB profile	22
B Experiment Tracking and Run Governance	22
C Operational TRL, Colab, and Evaluation Contract	23
D Curation Workflow and Publication Standard	24
E Reproducibility and Artifact Map	24
F Metric Interpretation	25

1 Introduction

Large language models are trained primarily by next-token prediction, while users ultimately care about properties such as helpfulness, relevance, correctness, clarity, and safety. RLHF addresses this mismatch by learning a reward signal from preference comparisons and then optimizing a language model against that signal [1, 4, 5]. In its canonical form, the procedure has three stages: supervised fine-tuning on preferred demonstrations, training a reward model on chosen/rejected pairs, and policy optimization with a KL constraint to a reference policy.

This project asks a deliberately modest but technically demanding question: can the complete procedure be made measurable and inspectable on a half-billion-parameter instruction model? The small policy makes iteration possible on rented notebook hardware, but it also makes the experiment a stress test of data quality, reward reliability, implementation detail, and evaluation discipline. The base model, QWEN2.5-0.5B-INSTRUCT, is already instruction-tuned and belongs to the Qwen2.5 family [8]. The work is therefore preference post-training rather than instruction tuning from an unaligned base.

The project also connects two settings that are often discussed separately. Earlier work in the same codebase implemented Trust Region Policy Optimization (TRPO), Natural Policy Gradient, and PPO for MuJoCo and Atari. TRPO constrains policy updates through a trust region [2]; PPO approximates the same conservative-update principle with a clipped surrogate objective [3]. In language-model RLHF, a second trust-region-like constraint appears through the KL penalty to a frozen SFT reference. The optimization machinery changes scale, but the central concern remains the same: improve expected reward without allowing destructive policy drift.

1.1 Contributions

The project makes five practical contributions:

1. It implements the full SFT–reward-model–PPO pipeline first with custom training code and then with TRL, preserving the custom path as a reproducible historical baseline.
2. It develops response-safe, long-context preprocessing for HELPS TEER3-PREFERENCE, including assistant-only loss masks, complete EOS-terminated completions, distinct PAD/EOS tokens, prompt-first truncation, and prompt deduplication for PPO.
3. It incorporates high-impact implementation details from the N+ reproduction of RLHF with PPO [6], including zero dropout, matched behavior probabilities, reward-model initialization from SFT, deliberate scalar-head initialization, reward centering, an RM-initialized critic, fixed-length generation, and strict handling of missing EOS.
4. It evaluates Base, SFT, and PPO on the same 2,017 prompts, stores full responses, supports resumable generation, and derives all pairwise tables from one shared response table.
5. It treats qualitative auditing as part of the result. Repetition, stopping, reward margins, domain effects, and reward-model mismatches are surfaced through deterministic triage, a balanced curated subset, and a public response explorer.

The source repository is available at <https://github.com/djdhillxn/rlhf>. The response explorer is available at <https://djdhillxn.github.io/projects/rlhf-comparison/>.

2 RLHF Formulation

2.1 Notation

Let x denote a formatted conversational prompt and $y = (y_1, \dots, y_T)$ an assistant response. The trainable policy is $\pi_\theta(y \mid x)$, the frozen SFT reference is $\pi_{\text{ref}}(y \mid x)$, and the scalar reward model is

$r_\phi(x, y)$. A response mask $m_t \in \{0, 1\}$ identifies generated assistant tokens so prompt and padding positions do not contribute to policy or value losses.

2.2 Supervised fine-tuning

SFT maximizes the likelihood of the preferred assistant response while masking the prompt:

$$\mathcal{L}_{\text{SFT}}(\theta) = - \sum_{t=1}^T m_t \log \pi_\theta(y_t \mid x, y_{<t}). \quad (1)$$

This is ordinary causal language modeling applied only to the completion. It serves two roles: it creates a policy adapted to the preference dataset, and it defines the reference distribution against which PPO drift is measured. Low-rank adaptation (LoRA) [9] updates small trainable matrices inside the attention and feed-forward projections while leaving the main backbone fixed.

2.3 Pairwise reward modeling

Each preference record contains a chosen response y^+ and a rejected response y^- for the same prompt. The reward model is trained with the Bradley–Terry logistic loss:

$$\mathcal{L}_{\text{RM}}(\phi) = - \log \sigma \left(r_\phi(x, y^+) - r_\phi(x, y^-) \right), \quad (2)$$

where σ is the sigmoid function. Pairwise accuracy is

$$\text{Acc}_{\text{RM}} = \frac{1}{N} \sum_{i=1}^N \mathbf{1} \left[r_\phi(x_i, y_i^+) > r_\phi(x_i, y_i^-) \right]. \quad (3)$$

The sign of an isolated score is not intrinsically meaningful: adding a constant to every reward leaves Equation 2 unchanged. This project therefore interprets margins and rankings rather than treating zero as a quality boundary. After training, a fixed demonstration-score offset is estimated and subtracted consistently during PPO and evaluation.

2.4 KL-regularized PPO

For a sampled response, the reward model supplies a terminal score while the reference policy supplies a token-level KL-shaped penalty:

$$r_t^{\text{KL}} = -\beta \left[\log \pi_\theta(y_t \mid x, y_{<t}) - \log \pi_{\text{ref}}(y_t \mid x, y_{<t}) \right], \quad (4)$$

$$R(x, y) = r_\phi(x, y) + \sum_{t=1}^T r_t^{\text{KL}}. \quad (5)$$

The coefficient β controls policy drift. Generalized advantage estimation (GAE) computes

$$\delta_t = r_t + \gamma V(s_{t+1}) - V(s_t), \quad (6)$$

$$\hat{A}_t = \delta_t + \gamma \lambda \hat{A}_{t+1}, \quad (7)$$

with $\gamma = 1$ and $\lambda = 0.95$ in the final run. Advantages are whitened over valid response tokens. For old-policy log probability $\log \pi_{\text{old}}$ and updated-policy log probability $\log \pi_{\theta}$, define

$$\rho_t(\theta) = \exp(\log \pi_{\theta}(y_t | s_t) - \log \pi_{\text{old}}(y_t | s_t)). \quad (8)$$

PPO minimizes the negative clipped surrogate:

$$\mathcal{L}_{\text{policy}} = -\mathbb{E}_t \left[\min \left(\rho_t \hat{A}_t, \text{clip}(\rho_t, 1 - \epsilon, 1 + \epsilon) \hat{A}_t \right) \right], \quad (9)$$

with $\epsilon = 0.2$. The critic uses a clipped squared-error objective, and the combined optimization loss is

$$\mathcal{L} = \mathcal{L}_{\text{policy}} + c_v \mathcal{L}_{\text{value}} - c_H \mathcal{H}(\pi_{\theta}). \quad (10)$$

The final configuration uses $c_v = 0.1$. These equations were implemented directly in the custom trainer and later delegated to the experimental TRL PPO trainer. The repository continues to own model preparation, sampling controls, EOS handling, reward centering, checkpoint metadata, and evaluation.

3 Model, Data, and Preprocessing

3.1 Base model and preference data

HELPSTEER3-PREFERENCE is a human-annotated preference dataset spanning general instruction following, code, STEM, and multilingual tasks [7]. Its breadth is useful for a general alignment study, but it also makes the task harder than short-form summarization: many preferred answers contain code, multi-step explanations, or long conversational context.

Table 1: Final data and model profile.

Item	Value
Base policy	QWEN2.5-0.5B-INSTRUCT
Approximate parameter count	0.5B
Raw training preference records	38,459
Raw validation preference records	2,017
Filtered SFT/RM training records	36,264
Filtered SFT/RM validation pairs	1,917
Unique PPO training prompts	24,018
Policy-suite evaluation prompts	2,017
Maximum SFT/RM sequence length	4,096 tokens
Maximum PPO/evaluation prompt length	3,072 tokens
PPO rollout response length	768 tokens
Evaluation generation cap	1,024 tokens

Invalid or tied preferences remove 2,195 training records and 100 validation records from SFT/RM training. PPO uses prompt-only records and removes 14,441 duplicate training prompts. The policy-suite evaluation is separate from this deduplicated PPO prompt cache and covers all 2,017 validation prompts.

3.2 Chat formatting and truncation

All stages use the Qwen chat template and the same tokenizer. The preprocessor constructs three views of each retained preference record:

- **SFT:** prompt plus chosen response, with a completion mask that excludes prompt tokens from Equation 1;
- **Reward model:** chosen and rejected sequences sharing the same prompt prefix;
- **PPO:** deduplicated, left-padded prompt token IDs.

PAD and EOS use distinct token IDs. Every SFT and reward-model completion ends in EOS. When a sequence exceeds the budget, the preprocessor removes the oldest non-system conversational turns first and left-truncates prompt tokens only as a final fallback. It does not cut a response merely to preserve more prompt history.

The project initially used a 1,024-token total limit. A direct length diagnostic showed that this discarded too much preference signal:

Table 2: Historical truncation diagnostic that motivated the 4,096-token training budget.

Limit	Train SFT	Train RM	Val. SFT	Val. RM
1,024	38.47%	40.82%	36.78%	39.49%
2,048	15.48%	16.47%	13.51%	14.87%
3,072	5.28%	5.83%	4.69%	5.32%
4,096	0.83%	1.00%	0.68%	0.89%

The final preparation report is consistent with this diagnostic: 304 of 36,264 SFT training prompts and 363 of 36,264 reward-model training prompts required truncation. The median chosen/SFT sequence length is 768 tokens and the 95th percentile is roughly 3,080 tokens. Long-context handling was therefore not cosmetic; it preserved a substantial part of the dataset.

4 Implementation and Experimental Infrastructure

4.1 From a custom trainer to TRL

The first implementation built the algorithm from low-level components: response masking, shifted token log-probabilities, token-level KL rewards, GAE, normalized advantages, clipped policy and value losses, adaptive KL control, LoRA checkpointing, and evaluation. This path was useful precisely because it failed in visible ways. Short response caps produced truncation; weak KL control or noisy adaptive estimates allowed drift; incorrect checkpoint paths invalidated comparisons; and high reward sometimes coincided with gibberish, empty outputs, or repeated text.

The final system uses Hugging Face TRL for trainer loops and standard serialization [10]. The migration was not a deletion of the project. Responsibilities were divided as follows:

Table 3: Ownership after migration to TRL.

Repository-owned	TRL-owned
HelpSteer3 parsing and filtering	SFT, reward, and PPO trainer loops
Qwen chat formatting and truncation	Mixed-precision/distributed preparation
Completion masks and prompt deduplication	Gradient accumulation and optimizer stepping
Model initialization and reward centering	PPO ratio, clipped objectives, and GAE
Fixed-length generation/EOS patches	Trainer checkpoint serialization
Run manifests and resolved configurations	Core training-loop metrics
Policy-suite evaluation and qualitative audit	PEFT trainer integration

The custom implementation remains frozen at Git commit:

6cbf214fcf1b91c7b756e303e533c2c86d2eba89

This preserves the learning history and provides a baseline for inspecting how trainer mechanics changed.

4.2 Direct preference optimization extension

The next controlled experiment uses Direct Preference Optimization (DPO) [12] from the same frozen SFT checkpoint. This comparison changes the alignment objective while holding the policy initialization, preference source, tokenizer, sequence budget, LoRA parameterization, and evaluation prompts fixed. It therefore tests whether direct offline preference learning is more robust than optimizing the learned scalar reward with online PPO at this model scale.

The DPO cache stores explicit **prompt**, **chosen**, and **rejected** fields. It applies the same response-safe 4,096-token preprocessing as SFT and reward modeling: complete EOS-terminated responses are preserved, old prompt turns are removed first, and only a final prompt-side token fallback is allowed. Tied pairs are excluded. The primary run uses one row per non-tied pair because HelpSteer3’s overall preference levels $\{-3, \dots, 3\}$ are ordinal labels rather than established linear utility distances. Preference magnitude, individual annotator scores, and directional agreement remain attached as metadata for stratified analysis. A separately named linear-replication run is available only as a sensitivity experiment.

The active DPO implementation uses the Hugging Face TRL `DPOTrainer`, a standard sigmoid loss with $\beta = 0.1$, zero dropout, LoRA rank 16, a 4,096-token prepared-data budget, one epoch, and effective batch size 32. Reference log probabilities are precomputed, and the reference policy is the same SFT model with the trainable adapter disabled. On Colab, training runs under `/content` while the SFT model, prepared DPO data, periodic complete Trainer checkpoints, and final artifacts are synchronized to Google Drive. Checkpoints are written every 200 optimizer updates and `resume_from_checkpoint=auto` restores optimizer, scheduler, RNG, and Trainer state after a runtime loss. The controlling notebook is `notebooks/rlhf_dpo_colab_pipeline.ipynb`; the canonical configuration is `configs/trl/qwen25_05b_helpsteer3_dpo.yaml`.

4.3 Implementation details adopted from N+

RLHF is unusually sensitive to details that can look incidental in pseudocode. The N+ reproduction [6], which itself reconstructs the summarization setup of Stiennon et al. [4], guided the final configuration:

Table 4: High-impact RLHF/PPO implementation details in the final path.

Detail	Implementation
Disable PPO dropout	PyTorch dropout probabilities and LoRA dropout are zero, preventing stochastic ratio noise between rollout and update.
Match behavior probabilities	Generation temperature is applied consistently when storing and recomputing log probabilities; checkpoint generation heuristics are neutralized.
Preserve completions and EOS	Prompt history is truncated before response tokens; SFT/RM completions end in EOS.
Initialize RM from SFT	The merged SFT model supplies the reward backbone.
Initialize reward head deliberately	The scalar head uses zero bias and standard deviation $1/\sqrt{d_{\text{hidden}} + 1}$.
Center rewards	Training uses a centering regularizer, then subtracts a persisted mean demonstration score.
Initialize critic from RM	The PPO value model begins from the trained reward-model weights.
Use a frozen SFT reference	PPO adapter weights are disabled for reference log probabilities.
Fixed-length EOS trick	Rollouts sample the configured length without stopping at EOS, are truncated at EOS afterward, and receive reward -1 if no EOS appears.
Whiten rewards/advantages	Reward whitening is enabled and TRL whitens masked advantages.
Adam numerical setting	Adam epsilon is 10^{-5} .

These choices reduce avoidable instability; they cannot repair a weak reward model, noisy preferences, or insufficient model capacity.

4.4 Reproducibility and Colab execution

Each stage writes a resolved YAML configuration, experiment manifest, summary, trainer metrics, and stage-specific artifacts. Smoke runs precede expensive runs. A “doctor” script validates the Python interpreter, CUDA visibility, trainer APIs, prepared datasets, tokenizer special tokens, and local/Drive write paths.

The Colab workflow keeps active datasets and checkpoints on local SSD for throughput and synchronizes periodic checkpoints and final artifacts to Google Drive. SFT and reward training support exact Transformers checkpoint resume. The experimental TRL PPO wrapper intentionally does not claim exact resume because the trainer does not expose a reliable optimizer/dataloader resume path in the used API. PPO continuation must instead be recorded as a new segment with explicit parent checkpoints.

Evaluation is resumable at the policy-output level. Each policy writes complete JSONL records as generation proceeds; “partial” denotes interrupt-safe stage output, not truncated response text. A fingerprint guards against reusing cached responses under a changed checkpoint or generation configuration.

5 Final Training Configuration and Results

5.1 Supervised fine-tuning

SFT trains LoRA modules in the attention projections (**q**, **k**, **v**, and **o**) and feed-forward projections (**gate**, **up**, and **down**). The effective batch size is $8 \times 4 = 32$ on one device.

Table 5: Final SFT configuration and outcome.

Setting or metric	Value
Trainer	TRL SFTTrainer
Epochs	1
LoRA rank / alpha	16 / 32
LoRA dropout	0
Effective batch size	32
Learning rate	5×10^{-6}
Scheduler	cosine, 3% warmup
Precision	bfloat16, TF32 enabled
Maximum total length	4,096
Train loss	1.0556
Validation loss	1.1127
Validation mean token accuracy	72.02%

“Mean token accuracy” is next-token accuracy over completion tokens. It is not response correctness, instruction-following accuracy, or a direct measure of human preference. The small decline from SFT to Base in the final reward-model comparison later in the report illustrates why token-level validation alone is not a sufficient checkpoint criterion.

5.2 Reward-model training

The reward model loads the merged SFT backbone, attaches a scalar sequence classification head, and trains LoRA rank 32. The first epoch was completed and synchronized; training then resumed from its saved checkpoint to two total epochs. The second-stage configuration uses an effective batch of 64 and centering coefficient 0.01. The persisted raw-score mean on 36,264 preferred demonstrations is -0.1985 (raw-score standard deviation 0.6594), and this offset is subtracted during PPO and evaluation.

Table 6: Final reward-model validation audit.

Metric	Value	Count
Overall pairwise accuracy	65.62%	1,917
Mean chosen-minus-rejected margin	0.3578	1,917
Final evaluation loss	0.6166	1,917
Code accuracy	70.23%	430
General accuracy	62.91%	914
STEM accuracy	59.26%	243
Multilingual accuracy	71.82%	330

Accuracy rises with preference strength: 61.23% for strength 1, 66.33% for strength 2, and 71.10% for strength 3. This is encouraging because clearer human preferences are easier to rank. The low STEM accuracy is the more consequential signal: a general reward model is least reliable in a domain where plausible but incorrect answers can be harmful.

The earlier custom reward model reached 71.62% pairwise validation accuracy, yet its downstream PPO policy did not beat Base. Conversely, the final TRL reward model has lower raw validation accuracy but participates in the best downstream run. This does not imply that lower accuracy is desirable. It shows that one aggregate RM metric cannot summarize initialization, reward scale, EOS handling, critic quality, policy optimization, and out-of-distribution policy outputs.

5.3 PPO alignment

The selected PPO run uses one rollout batch of 64 prompts per optimizer update (2 examples per device with 32 accumulation steps). Each rollout is reused for four PPO epochs. The run was configured for 12,000 episodes, corresponding to 188 planned updates, and was stopped for full evaluation after 6,400 episodes and 100 updates.

Table 7: Final PPO settings and observed training statistics.

Setting or metric	Value
Rollout response length	768 tokens
Rollout batch size	64 prompts
PPO epochs per rollout	4
Policy/value clip ranges	0.2 / 0.2
Value coefficient	0.1
γ / GAE λ	1.0 / 0.95
Sampling temperature / top- p	0.7 / 1.0
KL coefficient β	0.07
Missing-EOS reward	-1.0
Learning rate	3×10^{-6} , linear decay
Reward whitening	enabled
Evaluated episodes / updates	6,400 / 100
Mean response-level objective KL	1.8278
Final response-level objective KL	2.1648
Mean old-to-new approximate KL	0.000640
Mean PPO clip fraction	0.00557
Mean reward-model score in rollouts	-0.5898
Mean EOS count per 64 samples	44.57
Final EOS count per 64 samples	38

The logged `objective/kl` is a response-level sampled sum/aggregate, not a per-token KL target. It should not be compared directly with per-token values such as 0.01 or 0.05 without normalizing by valid response length. The much smaller `policy/approxkl_avg` measures the update from the rollout policy to the post-update policy, whereas `objective/kl` measures departure from the frozen SFT reference. Their different scales are expected.

Negative raw rollout scores are not evidence that PPO receives no learning signal. Reward scores are centered and can be negative for an entire batch; advantages depend on relative returns and value estimates. The final run remained numerically stable: update KL and clip fractions stayed small, the value loss fell from 25.32 on the first update to 5.50 on update 100, and no empty-response collapse appeared. Stability, however, does not rule out reward-model exploitation, which is why evaluation and auditing determine the final interpretation.

6 Experimental Progression

The final configuration was not chosen in one pass. Table 8 summarizes the main iterations without reproducing every debug run.

Table 8: Development path and principal lesson from each phase.

Phase	Experiment	Lesson
Short-context validation	96/128-token generations and small prompt budgets	Fast debugging exposed truncation, blank/EOS collapse, drift, and checkpoint path errors, but could not support credible qualitative evaluation.
Length diagnostic	1,024–4,096 total-token budgets	At 1,024 tokens, roughly 40% of SFT/RM examples truncate; 4,096 reduces this to about 1%.
Custom long-context SFT/RM	4,096-token SFT and two-epoch RM	The custom RM reached 71.62% pairwise accuracy, but judge accuracy alone did not ensure a useful PPO policy.
Custom PPO, 512 evaluation	397 of 400 updates, 512-token cap	Training did not collapse, but PPO won only 44.82% against Base and roughly 30% of outputs hit the cap.
Custom PPO, 1,024 evaluation	Same custom policy, larger generation allowance	Cap hits fell to 11.60% for PPO, while longer outputs exposed repetition, fabricated citations, and technical errors. The comparison was not a clean ablation because batch size also changed.
TRL smoke pipeline	Library-based SFT/RM/PPO with local-SSD and Drive synchronization	Validated stage interfaces, checkpoint loading, merged adapters, and evaluation before the expensive run.
Final TRL run	N+ details, two-epoch RM, 768-token PPO, full 1,024-token evaluation	Produced the first project run in which PPO narrowly beats Base under the learned reward model, while exposing a stronger repetition/stopping cost.

For context, the historical learned-reward comparisons were:

Table 9: Selected policy iterations. These are not controlled ablations: training code, reward models, decoding environments, and batch sizes differ.

Run	PPO vs Base	Mean delta	PPO cap hit	PPO repeat > 25%
Custom, eval cap 512	44.82%	−0.1327	29.70%	—
Custom, eval cap 1,024	38.92%	−0.2137	11.60%	16.11%
Final TRL, eval cap 1,024	50.92%	+0.6497	27.42%	31.88%

Three lessons shaped the final system. First, sequence budgets are part of the learning problem: short caps can hide both useful completions and late degeneration. Second, a reward model must be audited on policy-generated responses, not only held-out preference pairs. Third, small implementation differences – dropout, temperature-adjusted log probabilities, reward scale, EOS treatment, and critic initialization – are algorithmic choices in practice.

7 Final Policy-Suite Evaluation

7.1 Protocol

The evaluator generates one sampled response from each of Base, SFT, and PPO for every one of the 2,017 validation prompts. All policies use the same saved tokenizer, prompt formatting, 3,072-token prompt limit, 1,024-token generation limit, temperature 0.7, top- $p = 1.0$, and no repetition penalty or no-repeat n-gram constraint. The models load sequentially to manage memory. Batch size is 256, prompt order is fixed, and generation can resume from complete per-policy JSONL records.

Every prompt-response pair is scored with the two-epoch reward model. All overall winners, pairwise comparisons, domain tables, plots, and curation files derive from the same response table. This design prevents a pairwise comparison from accidentally using different generations. It does not turn the reward model into a human judge; it measures what the trained proxy prefers.

7.2 Overall results

Table 10: Three-way learned-reward evaluation on 2,017 prompts.

Policy	Wins	Win rate	Mean reward	Mean tok.	Median tok.	Cap hit
Base	718	35.60%	0.0803	398.0	331	8.82%
SFT	508	25.19%	0.0652	436.5	371	13.39%
PPO	775	38.42%	0.7300	577.8	520	27.42%
Tie	16	0.79%	—	—	—	—

Table 11: Pairwise comparisons under the learned reward model. “Right win rate” includes ties in the denominator.

Comparison	Left wins	Right wins	Ties	Right win rate	Mean Δ
Base vs SFT	1,158	837	22	41.50%	−0.0151
Base vs PPO	981	1,027	9	50.92%	+0.6497
SFT vs PPO	840	1,164	13	57.71%	+0.6648

PPO is the most frequent three-way winner and narrowly crosses 50% against Base. The median PPO-minus-Base reward delta is only +0.0254, while the mean is +0.6497. This gap indicates a skewed margin distribution: a subset of large positive deltas raises the mean. Because some extreme positive deltas are repetition failures, the aggregate mean cannot be interpreted without the audit in Section 8.




Figure 1: Reward-model score distributions for the three policies. PPO has a heavier positive tail, but distributional separation is limited and extreme scores require inspection.

7.3 Domain behavior

Table 12: PPO versus Base by prompt domain.

Domain	Prompts	PPO wins	Base wins	Ties	PPO win rate
Code	438	188	250	0	42.92%
General	931	529	399	3	56.82%
STEM	245	118	126	1	48.16%
Multilingual	403	192	206	5	47.64%

General prompts account for the overall advantage. PPO loses more often than it wins on code and remains slightly below Base on STEM and multilingual prompts. The domain pattern is consistent with the central limitation of a single generic reward model: it can learn broad stylistic and helpfulness preferences while remaining less reliable on correctness-sensitive tasks.

8 Qualitative Audit and Response Explorer

8.1 Stopping and repetition

The evaluation stores full response text, response length, EOS and cap status, reward margins, and word-level repeated 4-gram diagnostics. For a response with multiset of contiguous word 4-grams $\mathcal{G}$, the repeated fraction is

$$f_{\text{rep4}} = 1 - \frac{|\text{unique}(\mathcal{G})|}{|\mathcal{G}|}. \quad (11)$$

This is a triage feature, not a semantic quality metric. Code and technical prose can repeat legitimate terminology, while a loop can evade the metric by changing punctuation or one word.

Table 13: Stopping and repetition diagnostics.

Policy	Cap-hit rate	$f_{\text{rep4}} > 0.25$	$f_{\text{rep4}} > 0.50$
Base	8.82%	204 (10.11%)	61 (3.02%)
SFT	13.39%	405 (20.08%)	171 (8.48%)
PPO	27.42%	643 (31.88%)	319 (15.82%)




Figure 2: Response-token distributions. The large PPO mass at 1,024 tokens matches its 27.42% cap-hit rate and motivates explicit stopping analysis.

PPO’s longer median response can reflect useful additional detail, but it also creates more opportunity to repeat, hallucinate, or continue after the useful answer is complete. The cap spike in Figure 2 shows that the difference is not merely a smooth shift toward slightly longer answers.

8.2 Deterministic triage

Every evaluation row receives one mutually exclusive heuristic label based on reward deltas, EOS/cap behavior, and repetition:

Table 14: Full-set deterministic triage labels. These are navigation aids, not human or LLM-as-judge verdicts.

Label	Count	Meaning
Likely genuine PPO win	8	PPO beats Base and SFT by a strong reward margin, reaches EOS, avoids the cap, and stays below the repetition threshold.
Modest clean PPO win	354	PPO beats both baselines with low repetition and no cap hit, but with a smaller reward margin.
Needs manual review	924	No deterministic rule makes a strong call.
Repetition risk	228	PPO has elevated repeated 4-gram mass without meeting a more severe rule.
Reward-model false-positive risk	151	The reward model scores PPO highly despite repetition or cap warning signs.
Severe repetition failure	288	PPO reaches the cap with very high repeated 4-gram mass.
Strong PPO regression	64	PPO trails Base by a large reward margin.

Inspection of candidate extremes found both genuine local improvements and clear failures. Positive cases include more supportive responses, better coverage of multi-part instructions, and compact recoveries from capped or meandering baseline answers. Negative cases include repeated loops, prompt restatement, irrelevant continuation, malformed code, fabricated citations, and incorrect scientific procedures. Several high-reward failures are direct evidence of proxy mismatch: the policy discovered features the reward model liked without producing a response a careful reviewer should prefer. This is the practical form of reward-model overoptimization studied more generally by Gao et

al. [11].

8.3 Curated 100 and complete public artifact

The public JSON artifact retains all 2,017 full prompts and Base/SFT/PPO responses. Secret-shaped strings are redacted before export. A `curated` flag marks a balanced 100-example first-pass subset:

- 50 positive and 50 negative examples;
- 49 general, 21 multilingual, 19 code, and 11 STEM examples;
- the original reviewed 25 positive/25 negative set is retained;
- the expansion adds clean PPO gains by domain and balances the negative half across 15 reward-model mismatch risks, 15 severe repetition failures, 10 repetition risks, and 10 strong regressions.

The positive half contains all 8 likely-genuine heuristic wins, 41 modest clean wins, and one manually retained compact-recovery edge case whose automatic label remains “needs manual review.” Keeping both labels visible is intentional: curated polarity records why the example is informative, while heuristic triage records what the deterministic rule can actually support.

The response explorer loads the full artifact rather than a flattering demo file. Domain filters and the orthogonal Curated filter can be combined, and the page displays full prompt context, complete Base/PPO responses, reward scores, token counts, EOS/cap status, repetition metrics, and heuristic rationale. This makes the evaluation inspectable without requiring a notebook or local model weights.

9 Discussion

9.1 What the final run demonstrates

The strongest supported claim is a systems claim. The project successfully connects long-context data preparation, LoRA SFT, pairwise reward modeling, online PPO, checkpoint-safe execution, full-suite generation, quantitative comparison, and qualitative audit on a real, diverse preference dataset. PPO changes behavior measurably and produces useful local gains. Under the learned reward model, the selected policy narrowly beats Base overall and beats SFT by a larger margin.

The final run also demonstrates why stability is not alignment. Low update KL, small clip fractions, finite losses, and nonzero EOS counts show that optimization is behaving numerically. They do not show that the objective is correct. The repetition and cap metrics reveal a policy that can remain close enough to its reference for PPO diagnostics while still exploiting weaknesses in a scalar judge.

The custom-to-TRL transition was valuable for two different reasons. The custom implementation made algorithmic assumptions visible and produced the failures that motivated better safeguards. TRL then supplied a cleaner, better-tested training substrate for the final run. The project retained its distinct value because data contracts, model initialization, EOS behavior, experiment records, evaluation, and audit remained explicit rather than becoming an opaque library call.

9.2 Limitations

1. **The training judge is also the evaluation judge.** The same reward model supplies PPO rewards and final policy scores. A policy can improve this metric by exploiting judge-specific blind spots. The evaluation is therefore a reward-model-based comparison, not an independent human preference result.
2. **No blinded human study is reported.** Manual review informed curation and interpretation, but it was not a preregistered, multi-rater, label-blind preference study with agreement statistics.
3. **One generation per policy is evaluated.** Sampled decoding at temperature 0.7 has variance. Repeated seeded samples and paired bootstrap confidence intervals would produce stronger

uncertainty estimates.

4. **Response length is a confound.** PPO is substantially longer. Longer answers can be more complete, but can also accumulate judge-preferred formatting, unsupported claims, and repetition.
5. **The model and reward model are small.** A 0.5B policy is economical and debuggable, but capacity limits instruction following, factuality, long-context coherence, and value estimation.
6. **The historical comparisons are not clean ablations.** The 512- and 1,024-token custom evaluations differed in batch size and software details, so they document operational learning rather than isolate one causal variable.
7. **The final checkpoint is practical rather than statistically optimal.** Training stopped after 100 updates for evaluation under a deadline. No independent human metric selected it from a checkpoint sweep.
8. **Heuristic auditing is incomplete.** Word-level repetition and a small sensitive-term lexicon can prioritize review but cannot establish factuality, safety, or harmlessness.

These limitations are not footnotes to an otherwise definitive win. They define the boundary of the result.

10 Future Work

The next phase should improve the evidence and evaluator before simply increasing PPO compute. The agenda below preserves the complete research plan that emerged from the custom and TRL experiments.

10.1 Establish a controlled evaluation protocol

Future comparisons should make one intervention attributable at a time. The historical 512- and 1,024-token suites changed batch size as well as response budget, and most generations changed even when the shorter response had not reached its cap. A controlled token-budget study should hold the checkpoint, software environment, precision, batch size, prompt order, decoding settings, seed, and evaluator implementation fixed while testing 256, 512, 768, and 1,024 new tokens. It should also record whether a longer response is an exact continuation of its shorter counterpart.

Greedy decoding supplies a reproducible trajectory but not a distributional evaluation. Repeated sampled generations with fixed seeds should measure both mean quality and response variance. Because every policy answers the same prompts, overall and domain-level confidence intervals should use paired bootstrap resampling.

Blinded human review should become a first-class signal. Prompts should be stratified by domain, learned-reward margin, response length, repetition risk, and curated polarity. Reviewers should not see policy names or reward scores and should separately rate correctness, relevance, completeness, clarity, harmlessness, and unnecessary verbosity. Inter-rater agreement and adjudication rates would distinguish judge error from genuinely ambiguous preferences.

Length must be controlled explicitly. Reports should compare reward within response-length buckets, win rates among similarly sized answers, and reward delta against token delta. This separates substantive improvement from a judge that rewards formatting or verbosity.

10.2 Improve reward-model reliability

The final reward model reaches 65.62% held-out pairwise accuracy, and several extreme scores still contradict visible response quality. Its next dataset should include hard negatives from the policy distribution: repeated loops, prompt restatements, irrelevant continuations, malformed code, fabricated citations, incorrect STEM procedures, language drift, and answers whose useful prefix

degrades later. These examples are more relevant to policy optimization than random rejected HelpSteer3 responses and should also form a held-out stress test.

One scalar preference target should not be assumed to represent every desired property. A stronger evaluator could retain a general preference score while adding auxiliary heads or separate judges for relevance, factuality, repetition, safety, style, code execution, and language consistency. Deterministic checks for empty text, missing EOS, repeated n -grams, language mismatch, or failure to answer a direct question can guard known blind spots, although they must not replace learned or human preference.

Reward uncertainty should be measured. Ensembles trained from different seeds or splits can identify unstable decisions; policy rewards can be discounted where judges disagree, and evaluation can separate high- and low-confidence comparisons. Calibration should test whether larger reward margins actually correspond to more reliable preferences. Domain rebalancing is particularly important because STEM accuracy is materially lower than code and multilingual accuracy. Native-speaker review is also needed for multilingual cases.

Long-response examples deserve an explicit curriculum. Batches should balance short, medium, and long answers and include pairs whose openings are similar but whose later continuations differ sharply in quality. This teaches the judge that a strong prefix does not excuse a repetitive or incorrect ending. Scaling the reward model before the policy may be safer: a larger policy can exploit a weak judge more effectively.

10.3 Strengthen supervised fine-tuning

SFT determines both policy initialization and the frozen reference. Future data should therefore be filtered for factuality, coherent stopping, and relevance rather than treating the preferred label as automatic proof of quality. Duplicate or near-duplicate responses, broken markup, unsupported citations, and repeated text should be removed or repaired. A smaller clean set may be preferable to a larger noisy one.

Data composition should match intended use. Domain-stratified sampling and loss monitoring would prevent general English prompts from dominating. Code examples should be executable and unit tested; STEM examples should contain verified derivations and grounded citations; multilingual examples should be natural rather than translation-only. A sequence-length curriculum could teach instruction following and stopping on shorter high-quality responses before introducing long-context examples. Any packing strategy must preserve response-only masks and avoid confusing mixtures of unrelated conversations.

SFT checkpoint selection should compare epochs instead of assuming the final checkpoint is best. Preferred-response perplexity is insufficient; selection should include response-level preference, repetition, stopping, factuality, and downstream alignment behavior.

10.4 Redesign PPO around observed failures

A direct 1,024-token PPO run is useful only as a controlled experiment. It should start from the same SFT checkpoint and compare 512-, 768-, and 1,024-token policies at matched update-token budgets. A curriculum from 256 or 512 tokens to longer horizons may be more stable than beginning with the largest allowance. Missing-EOS rewards, length-aware terminal rewards, and repetition penalties should be introduced gradually and monitored for short-answer collapse.

The reward should distinguish a useful prefix from a degraded suffix. A single terminal score gives every response token the same sequence-level outcome. Process or segment rewards, or at least scoring the full answer and several prefixes, can reveal when continuing became harmful. Late-token KL diagnostics and an adaptive target-KL controller could likewise show whether repetition begins

after the policy departs from its reference late in the answer.

Prompt difficulty affects raw reward, so normalization across unrelated prompts can obscure the signal. Several candidates per prompt with within-prompt centering provide a cleaner comparison; RLOO- or GRPO-style baselines are relevant here. Checkpoint selection must use learned reward, KL, EOS rate, empty rate, cap-hit rate, repetition, length, human preference, and domain correctness together. Training should stop when reward improves while these safety or quality measures deteriorate.

10.5 Compare alternative preference objectives

PPO remains valuable because it connects the work to trust-region policy optimization, but it should not be the only alignment baseline. The controlled DPO extension uses the same SFT checkpoint and preference pairs while avoiding online rollouts, a learned reward function in the update loop, and a value model. IPO, KTO, ORPO, and related objectives provide additional robustness and data-requirement comparisons. Together they can help isolate whether the bottleneck is reward modeling, value estimation, online RL, model capacity, or the preference data.

Offline methods cannot automatically learn from new policy-specific failures. A later online loop could alternate generation, human or high-quality judge labeling, reward-model updates, and policy updates, turning the static dataset into active learning around the policy’s observed errors.

10.6 Improve decoding and stopping

Inference settings are part of policy behavior. The final suite uses sampled decoding at temperature 0.7 with no repetition penalty or no-repeat constraint. Deployment-oriented evaluation should compare greedy decoding, temperature and top- p sweeps, modest repetition penalties, no-repeat constraints, and task-aware stopping. These controls can suppress visible degeneration but may also damage legitimate repetition in code or structured prose.

Dynamic token budgets are preferable to one universal cap. Classification, extraction, and title prompts need little space, while code explanations and long-form synthesis may need much more. A task classifier or length predictor can allocate budgets while reports measure both quality and compute. Stopping quality itself should be measured through EOS rate, cap-hit rate, late-response repetition, and reward changes between full answers and shorter prefixes.

10.7 Add factuality, execution, and safety evaluators

Correctness-sensitive tasks admit stronger checks than a generic preference score. Code can be parsed, linted, and executed in isolated tests; SQL can run against temporary schemas; shell snippets can receive static analysis. STEM evaluation should check equations, units, factual claims, and citation existence, using retrieval or specialist review where appropriate. Abstention or calibrated uncertainty should be preferable to unsupported confidence.

Safety should use structured suites for self-harm, medical advice, harassment, sexual content, privacy, and dangerous procedural assistance. False refusals must be measured alongside unsafe compliance because refusing benign prompts is also an alignment failure.

10.8 Scale model and systems capacity deliberately

Repeating the controlled pipeline at 1.5B or 3B scale would test whether stronger base capability turns the same preference signal into more reliable gains. Data and evaluation must remain fixed across scales. Larger runs would need sharding, activation checkpointing, efficient attention,

packed sequences, quantized frozen reference/reward models, inference-oriented rollout engines, and potentially asynchronous generation.

Policy scale and reward-model scale should be separate experimental axes. Improving the judge may benefit training and evaluation simultaneously, whereas scaling only the policy may amplify exploitation.

10.9 Reproducibility and artifact governance

Every evaluation should retain its resolved configuration, package versions, source revision, checkpoint hashes, prompt-order hash, generation signature, and environment. Resume fingerprints should prevent stale cached responses from being reused under changed settings. Reports should be generated from lightweight commands, raw response tables should remain immutable, and new runs should use versioned directories rather than overwrite evidence.

10.10 Practical next sequence

The highest-value next step is a blinded human review of a stratified subset, paired with a controlled 512/768/1,024-token evaluation under identical conditions. The second step is retraining and stress-testing the reward model on hard negatives from the audit. The third is a comparison among conservative PPO, length-curriculum PPO, and the DPO baseline, with checkpoint selection based on human preference and the full diagnostic set rather than reward alone. This sequence targets observed bottlenecks and is more informative than merely extending the current PPO run.

11 Conclusion

This project completes a three-stage RLHF pipeline for a small instruction model and, more importantly, makes the pipeline inspectable. Long-context SFT preserves nearly all preference responses. The reward model learns a usable but imperfect preference signal. PPO remains stable and produces the first run in the project that narrowly exceeds Base under that learned signal. The same evaluation shows that the policy is longer, more frequently capped, and far more repetitive.

The right conclusion is therefore neither “PPO failed” nor “PPO solved alignment.” The project demonstrates that an individual practitioner can build and study the complete RLHF stack, reproduce many subtle implementation details, diagnose reward hacking, and publish the evidence needed to inspect both successes and failures. Its most durable result is the combination of training, evaluation, and honest audit. Future progress depends less on running PPO for a larger number of episodes than on improving the judge, controlling stopping, and validating preferences independently.

A Detailed Historical Experiment Ledger

This appendix retains the operational facts behind the compact progression in Table 8. Historical paths and configurations are evidence of what was run; they are not active recommendations.

A.1 Initial validation and sequence-length diagnosis

Early 96- and 128-token generations accelerated debugging but often stopped inside code blocks or explanations. These runs exposed multilingual or gibberish drift, vulgar debug-pattern text, EOS and blank-response collapse, overaggressive reward shaping, and invalid comparisons caused

by incorrect checkpoint paths. The resulting safeguards included prompt sanitation, KL anchoring, empty-response checks, repeated-text diagnostics, and explicit checkpoint validation.

The sequence-length diagnostic established why the project moved from 1,024 to 4,096 total tokens:

Table 15: Fraction of records requiring truncation at each total-token limit.

Limit	Train SFT	Train RM	Valid. SFT	Valid. RM
1,024	38.47%	40.82%	36.78%	39.49%
2,048	15.48%	16.47%	13.51%	14.87%
3,072	5.28%	5.83%	4.69%	5.32%
4,096	0.83%	1.00%	0.68%	0.89%

A.2 Frozen custom SFT, reward, and PPO baseline

The pre-TRL lineage is frozen at commit `6cbf214fcf1b91c7b756e303e533c2c86d2eba89`. Custom SFT used QWEN2.5-0.5B-INSTRUCT, 4,096 total tokens, two epochs, batch size 6, accumulation 3, learning rate 5×10^{-6} , and LoRA rank 16. It completed 4,030 optimizer steps and wrote to `outputs/rlhf/qwen25_05b_helpsteer3_sft_4096/`.

The two-epoch custom reward model wrote to `outputs/rlhf/qwen25_05b_helpsteer3_reward_4096_epoch2/` and produced the following held-out audit:

Table 16: Historical custom reward-model result.

Metric	Value
Validation pairs	1,917
Accuracy / loss	71.62% / 0.9734
Mean chosen–rejected margin	0.9094
Code / General accuracy	74.88% / 71.01%
STEM / Multilingual accuracy	63.37% / 75.15%

The custom long-512 PPO run used a 3,072-token prompt budget, 512-token rollouts, learning rate 3×10^{-7} , policy clip 0.06, one PPO epoch, rollout batch 2, and an adaptive KL coefficient initialized at 0.18 with a 0.14 minimum. It completed 397 of 400 requested updates and ended at KL coefficient 0.14. Rewards were clipped to $[-5, 5]$; scalar reward whitening was absent; length penalty was 0.002; missing-EOS penalty was zero; responses shorter than 32 tokens received a 0.5 penalty; group-relative normalization was disabled. Generation applied repetition penalty 1.05 and blocked repeated 4-grams. Prompt batches were left padded, and response masks excluded padding from PPO losses. Its output is `outputs/rlhf/qwen25_05b_helpsteer3_ppo_4096_epoch2_long512/`.

A.3 Historical policy-suite results

The 512-token suite used all 2,017 validation prompts. Three-way wins were Base 827 (41.00%), SFT 556 (27.57%), PPO 525 (26.03%), and 109 ties (5.40%). Mean rewards were -3.5339 , -3.4280 , and -3.6666 ; median response lengths were 332, 363, and 356 tokens; and cap-hit rates were 29.90%, 29.95%, and 29.70%. Its complete pairwise table is retained below.

Table 17: Historical custom 512-token pairwise evaluation.

Comparison	Left	Right	Ties	Right rate	Mean Δ
Base vs SFT	1,044	927	46	45.96%	+0.1059
Base vs PPO	1,068	904	45	44.82%	−0.1327
SFT vs PPO	965	898	154	44.52%	−0.2386

The later 1,024-token suite produced three-way wins of Base 978 (48.49%), SFT 475 (23.55%), PPO 467 (23.15%), and 97 ties (4.81%). Mean rewards were -3.3634 , -3.6114 , and -3.5771 ; median lengths were 334, 360, and 363; cap-hit rates were 8.08%, 10.16%, and 11.60%; empty rates were zero.

Table 18: Historical custom 1,024-token pairwise evaluation.

Comparison	Left	Right	Ties	Right rate	Mean Δ
Base vs SFT	1,215	763	39	37.83%	−0.2480
Base vs PPO	1,190	785	42	38.92%	−0.2137
SFT vs PPO	963	892	162	44.22%	+0.0343

PPO’s 785 Base wins had mean margin $+1.6210$, while its 1,190 losses had mean margin -1.4315 . Against SFT, its 892 wins averaged $+1.3503$ and its 963 losses -1.1789 . Domain PPO win rates against Base were 26.26% for code (115/438), 44.47% for general (414 wins, 487 losses, 30 ties), 39.59% for STEM (97/245 with two ties), and 39.45% for multilingual (159 wins, 236 losses, eight ties).

Increasing the evaluation cap reduced Base/SFT/PPO cap hits from 603/604/599 to 163/205/234. It also exposed more PPO repetition: 224 responses exceeded 25% repeated 4-grams in the 512-token suite versus 325 (16.11%) in the 1,024-token suite; severe cases above 50% rose from 74 to 156. Manual review found fabricated citations, incorrect chemistry, irrelevant continuations, and high-reward loops. This was not a clean token-cap ablation: batch size changed from 8 to 128, and bfloat16 batching, environment, or evaluator revisions could change close greedy continuations.

A.4 Historical A100 40GB profile

An earlier speed-oriented Colab profile used full bfloat16 weights rather than 4-bit loading. Reward training used batch 8, accumulation 2 (effective 16), all filtered HelpSteer3 pairs, and refreshed token throughput, CUDA memory, CSV files, and plots every configured artifact interval. Batch 12 was the suggested increase below roughly 28–30GB peak memory; batch 6 was the fallback near 36GB or after OOM.

Historical PPO used full bfloat16 policy/reference/reward models, rollout batch 16, minibatch 4, prompt length 768, response length 160, checkpoints every 50 updates, and metric refresh every 25 updates. Rollout batch 24 was the first recommended utilization increase below 30–32GB; minibatch 6 or 8 was reserved for runs with measured headroom. The early evaluator covered 200 prompts; the completed suite superseded it with all 2,017.

B Experiment Tracking and Run Governance

The project uses a filesystem contract rather than a tracking server. Every new run directory contains the following where applicable:

Table 19: Run-tracking contract.

Artifact	Purpose
<code>config_resolved.yaml</code>	Exact configuration consumed by the run.
<code>experiment_manifest.json</code>	Identity, intent, config hash, command, source state, environment, attempts, status, and summary.
<code>run_metadata.json</code>	Backward-compatible hardware/software record.
<code>*_metrics.jsonl</code> / CSV	Step-level measurements.
Run or evaluation summary	Final aggregate result.
Checkpoints, samples, plots	Stage-specific evidence.

An interrupted process leaves its manifest in **running** state. Reusing the directory appends an attempt and records whether the resolved configuration changed; successful completion updates status and summary. Configs can declare non-behavioral intent—experiment ID, name, group, tags, hypothesis, and parent runs—while algorithmic choices remain in functional sections such as generation, reward shaping, KL, and PPO.

Runs are compared with `python -m scripts.rlhf_compare_runs RUN_A RUN_B`; metadata-only fields are hidden by default. The command supports `-include-metadata`, `-format json`, and `-output comparison.md`. The frozen custom baseline remains under `experiments/baselines/qwen25_05b_helpsteer3_ppo_long512/`; its YAML record is authoritative after removal of its redundant prose README.

The final lightweight records live under `rlhf_runs_lightweight_export/qwen25_05b_helpsteer3_trl_a100_full/full/`. Stage sources of truth are SFT `run_summary.json`, reward epoch-1 and epoch-2 summaries/audits, PPO checkpoint-100 `trainer_state.json`, the full policy-suite summary and sample table, and the deep-curation artifacts. The PPO directory name retains an obsolete `r512` token; the executed notebook override of 768 tokens takes precedence.

C Operational TRL, Colab, and Evaluation Contract

The stage datasets share one filtered source but have distinct contracts: SFT stores token IDs and a completion mask; reward modeling stores paired chosen and rejected token IDs with a common prefix; PPO stores deduplicated left-padded prompts; DPO stores explicit prompt/chosen/rejected text plus preference-strength and annotator metadata. Prompt history is shortened before any final prompt-side truncation, responses retain EOS, and PAD is distinct from EOS.

The preflight doctor validates the exact interpreter, package versions, CUDA, trainer APIs, tokenizer special tokens, prepared data, and local/Drive write permissions before expensive work. Evaluation always loads the SFT tokenizer; Base is resized once for its distinct PAD token. Neutral generation controls are explicit so checkpoint defaults cannot create asymmetric decoding. Policy list overrides accept either dotted or bracketed indices. The policy labeled Base must keep a null checkpoint and load the configured model name; pointing it at SFT would silently make the Base and SFT columns identical.

Active work uses Colab local SSD for model shards, datasets, logs, and Trainer state, with periodic checkpoints and final artifacts copied to mounted Drive. This is faster and keeps model tensors outside Git. SFT, reward, and DPO support exact Trainer resume. DPO’s `auto` mode restores model, optimizer, scheduler, RNG, and Trainer state. The experimental PPO API used in this project does not expose trustworthy exact resume; a continuation is a new segment with an explicit parent, the original SFT reference, and separately recorded optimizer history.

Evaluation is resumable per policy. Its “partial policy outputs” are complete JSONL rows written incrementally; *partial* describes interrupt-safe run completion, not shortened prompts or responses.

A fingerprint prevents cached rows from being reused when checkpoint or generation settings change. TRL PPO APIs are experimental, and the repository intentionally validates the installed stack instead of assuming that an unpinned upgrade is compatible. Every dependency change requires a smoke run and a recorded environment before an expensive experiment.

D Curation Workflow and Publication Standard

The audit scans every full Base/SFT/PPO response for length, cap and EOS status, empty text, learned-reward margins, repeated word-level 4-grams, and a small sensitive-term lexicon. It writes candidate tables for low-repetition wins, strong regressions, repetition risk, and high-reward mismatch risk. Neither the repetition statistic nor lexical scan is a semantic or safety classifier.

Table 20: Curation artifacts retained by the final workflow.

Artifact	Purpose
<code>policy_suite_samples.jsonl</code> / CSV	Full prompts, responses, and quantitative diagnostics.
<code>qualitative_audit_auto.md</code>	Generated aggregate audit and candidate summary.
<code>curation_qualified_ppo_candidates.csv</code>	Low-repetition PPO wins over both comparison policies.
<code>curation_strong_ppo_losses.csv</code>	Large learned-reward regressions.
<code>curation_ppo_repetition_risks.csv</code>	Responses above the repeated 4-gram threshold.
<code>curation_reward_model_mismatches.csv</code>	High reward combined with repetition or stopping warnings.
<code>ckpt100_judged_examples.csv</code>	Deterministic labels for all 2,017 rows.
<code>ckpt100_deep_curation_pools.json</code>	Candidate pools used for manual selection.
<code>portfolio_full_policy_comparisons_final_public.json</code>	Delayed full response-explorer payload.
<code>portfolio_curated_100_manifest_finaltrl.json</code>	Balanced public first-pass subset.

The deterministic labels in Table 14 are produced only from recorded diagnostics and metadata. They are not manual judgments. Review must cover clean wins, modest or near-tie wins, strong losses, repetition, proxy mismatches, correctness-sensitive code/STEM failures, and multilingual drift. The first public curation contained 50 examples: 8 likely genuine wins, 16 modest clean wins, 9 reward-model false-positive risks, 9 severe repetition failures, 7 strong regressions, and one manual-review edge case. It was later expanded to the balanced 100-example subset described in the main audit, while the complete 2,017-row artifact remained public.

Before calling an example an improvement, a reviewer must check instruction following, correctness, completeness, relevance, repetition, safety, and whether the reward advantage reflects substance rather than length or format. Code should be executed when possible; technical claims and citations require independent verification. Publication must show both gains and failures. The defensible conclusion is that PPO narrowly leads under its learned judge while being longer, more capped, and more repetitive—not that PPO universally improves the model.

E Reproducibility and Artifact Map

Table 21: Primary project artifacts used in this report.

Artifact	Purpose
Project overview	<code>README.md</code> . Current project overview and final headline metrics.
Academic report	<code>docs/rlhf_project_report.tex</code> and PDF. Concise scientific account of the method, final results, interpretation, and controlled DPO extension.

Artifact	Purpose
Technical record	<code>docs/rlhf_technical_companion.tex</code> and PDF. Chronological experiments, trainer migration, implementation details, Colab/checkpoint contract, audit, artifact governance, and complete research agenda.
Colab reference	<code>docs/rlhf_colab_training.ipynb</code> . Historical executed notebook retained as an operational record; the active DPO driver is under <code>notebooks/</code> .
Prepared data	<code>rlhf_runs_lightweight_export/.../data/preparation_report.json</code> . Filtered counts, truncation counts, and sequence-length distributions.
SFT summary	<code>rlhf_runs_lightweight_export/.../sft/run_summary.json</code> . Final SFT training summary.
Reward audit	<code>rlhf_runs_lightweight_export/.../reward_epoch2/reward_audit.json</code> . Final reward-model accuracy, margins, and domain/language breakdown.
PPO trainer state	<code>rlhf_runs_lightweight_export/.../checkpoint-100/trainer_state.json</code> . Selected PPO training history through 100 updates.
Policy-suite summary	<code>rlhf_runs/checkpoints_ckpt100_full/policy_suite_summary.json</code> . Authoritative final policy and pairwise evaluation.
Full response artifact	<code>rlhf_runs/portfolio_full_policy_comparisons_final_trl.json</code> . All 2,017 full prompts and policy outputs, heuristic labels, and curated flags.
Curated manifest	<code>rlhf_runs/portfolio_curated_100_manifest_final_trl.json</code> . Compact manifest of the balanced curated subset.

The active command sequence is:

```
python -m scripts.rlhf_trl_prepare_data \
  --config configs/trl/qwen25_05b_helpsteer3_sft.yaml
python -m scripts.rlhf_trl_train_sft \
  --config configs/trl/qwen25_05b_helpsteer3_sft.yaml
python -m scripts.rlhf_trl_train_reward_model \
  --config configs/trl/qwen25_05b_helpsteer3_reward.yaml
python -m scripts.rlhf_trl_train_ppo \
  --config configs/trl/qwen25_05b_helpsteer3_ppo.yaml
python -m scripts.rlhf_evaluate_policy_suite \
  --config configs/trl/qwen25_05b_helpsteer3_eval_suite.yaml
python -m scripts.rlhf_trl_prepare_data \
  --config configs/trl/qwen25_05b_helpsteer3_dpo.yaml
python -m scripts.rlhf_trl_doctor \
  --config configs/trl/qwen25_05b_helpsteer3_dpo.yaml --stage dpo
python -m scripts.rlhf_trl_train_dpo \
  --config configs/trl/qwen25_05b_helpsteer3_dpo.yaml
python -m scripts.rlhf_audit_policy_suite \
  --eval-dir rlhf_runs/checkpoints_ckpt100_full \
  --base-label base --sft-label sft_trl \
  --ppo-label ppo_exact_ckpt100
python -m scripts.rlhf_sync_runs --mode lightweight
```

The executed Colab notebook is the source of truth when a directory name or base YAML conflicts with a runtime override. In particular, the PPO directory retains an older `r512` label, while the selected run used 768-token rollouts and KL coefficient 0.07.

F Metric Interpretation

SFT mean token accuracy

Fraction of preferred completion tokens predicted correctly under teacher forcing; not response-level task accuracy.

Reward-model pairwise accuracy

Fraction of held-out pairs for which the chosen response scores above the rejected response.

Policy win rate

Fraction of evaluation prompts for which one policy receives the higher learned reward. It is a proxy-judge result, not a human preference rate.

Objective KL

Sampled response-level departure from the frozen SFT reference. It is not the same statistic as PPO’s old-to-new update approximate KL.

Cap hit

The response uses the full evaluation generation budget and does not stop naturally before 1,024 new tokens.

Repeated 4-gram fraction

Fraction of word-level 4-gram positions attributable to repeated 4-grams, used only to prioritize qualitative review.

References

- [1] P. F. Christiano, J. Leike, T. Brown, M. Martic, S. Legg, and D. Amodei. Deep reinforcement learning from human preferences. In *Advances in Neural Information Processing Systems*, 2017. <https://arxiv.org/abs/1706.03741>.
- [2] J. Schulman, S. Levine, P. Abbeel, M. Jordan, and P. Moritz. Trust Region Policy Optimization. In *Proceedings of the 32nd International Conference on Machine Learning*, 2015. <https://arxiv.org/abs/1502.05477>.
- [3] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov. Proximal Policy Optimization Algorithms. arXiv:1707.06347, 2017. <https://arxiv.org/abs/1707.06347>.
- [4] N. Stiennon, L. Ouyang, J. Wu, D. Ziegler, R. Lowe, C. Voss, A. Radford, D. Amodei, and P. F. Christiano. Learning to summarize with human feedback. In *Advances in Neural Information Processing Systems*, 2020. <https://arxiv.org/abs/2009.01325>.
- [5] L. Ouyang et al. Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems*, 2022. <https://arxiv.org/abs/2203.02155>.
- [6] S. Huang, M. Noukhovitch, A. Hosseini, K. Rasul, W. Wang, and L. Tunstall. The N+ Implementation Details of RLHF with PPO: A Case Study on TL;DR Summarization. arXiv:2403.17031, 2024. <https://arxiv.org/abs/2403.17031>.
- [7] Z. Wang, J. Zeng, O. Delalleau, H.-C. Shin, F. Soares, A. Bukharin, E. Evans, Y. Dong, and O. Kuchaiev. HelpSteer3-Preference: Open Human-Annotated Preference Data across Diverse Tasks and Languages. arXiv:2505.11475, 2025. <https://arxiv.org/abs/2505.11475>.
- [8] Qwen Team. Qwen2.5 Technical Report. arXiv:2412.15115, 2024. <https://arxiv.org/abs/2412.15115>.
- [9] E. J. Hu et al. LoRA: Low-Rank Adaptation of Large Language Models. In *International Conference on Learning Representations*, 2022. <https://arxiv.org/abs/2106.09685>.
- [10] L. von Werra et al. TRL: Transformer Reinforcement Learning. Hugging Face, 2020–2026. <https://github.com/huggingface/trl>; https://huggingface.co/docs/trl/ppo_trainer.
- [11] L. Gao, J. Schulman, and J. Hilton. Scaling Laws for Reward Model Overoptimization. In *Proceedings of the 40th International Conference on Machine Learning*, 2023. <https://arxiv.org/abs/2210.10760>.

-
- [12] R. Rafailov, A. Sharma, E. Mitchell, C. D. Manning, S. Ermon, and C. Finn. Direct Preference Optimization: Your Language Model is Secretly a Reward Model. In *Advances in Neural Information Processing Systems*, 2023. <https://arxiv.org/abs/2305.18290>.